

# High Performance Transaction Processing for NVRAM

Steven Pelley, Thomas F. Wenisch  
University of Michigan

# Nonvolatile memory (NVRAM)

- New NVRAMs provide fast durable storage (phase change, memristor, STT-RAM)

Storage technology	Random read latency	Durable?
Disk	10ms	✓
Flash	90μs	✓
DRAM	100ns	✗
NVRAM	50-1000ns [IBM]	✓

*NVRAM enables **persistent** main memory*

# Recoverable systems

- Must protect data against server failure
  - E.g., file systems and databases
  - Typically rely on disk, flash for persistent storage
- NVRAM accelerates persistent storage access
  - Byte-addressable (native memory instructions) further decrease execution overheads
- Create recoverable, high performance data structures (main-memory/DRAM speed)

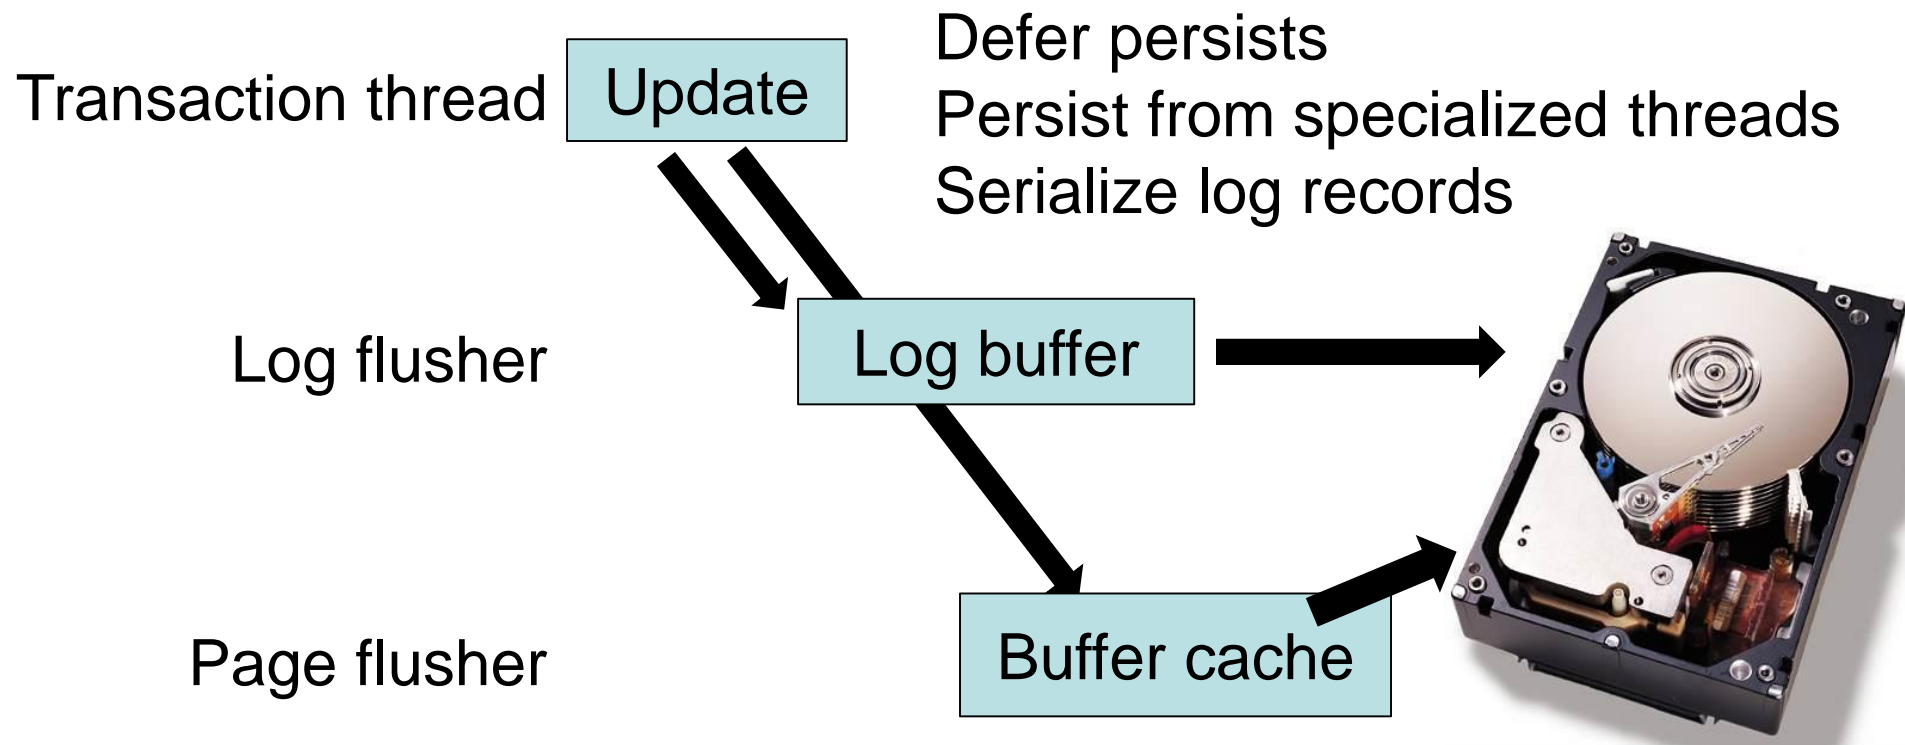
*NVRAM requires new recovery mechanisms*

# Optimizing NVRAM

- Case study of transaction processing [VLDB]
  - NVRAM as a disk replacement (ARIES/WAL)
  - Fast NVRAM imposes frequent persist ordering
  - New software design to avoid ordering delays
- Future work: explore desirable optimizations and semantics for NVRAM memory systems

# NVRAM as a disk replacement

## Write Ahead Logging (WAL) via ARIES




*Complexity necessary for slow disk now excessive overhead*

# Alternative: order persists immediately

- Persist barriers [BPFS: SOSP '09]
  - Persists following barrier instruction may not complete before persists prior to barrier
  - Persists that are not separated by barrier instruction are *concurrent*
- Barriers delay execution
  - Frequency of barriers and *exposed barrier latency* determine resulting performance

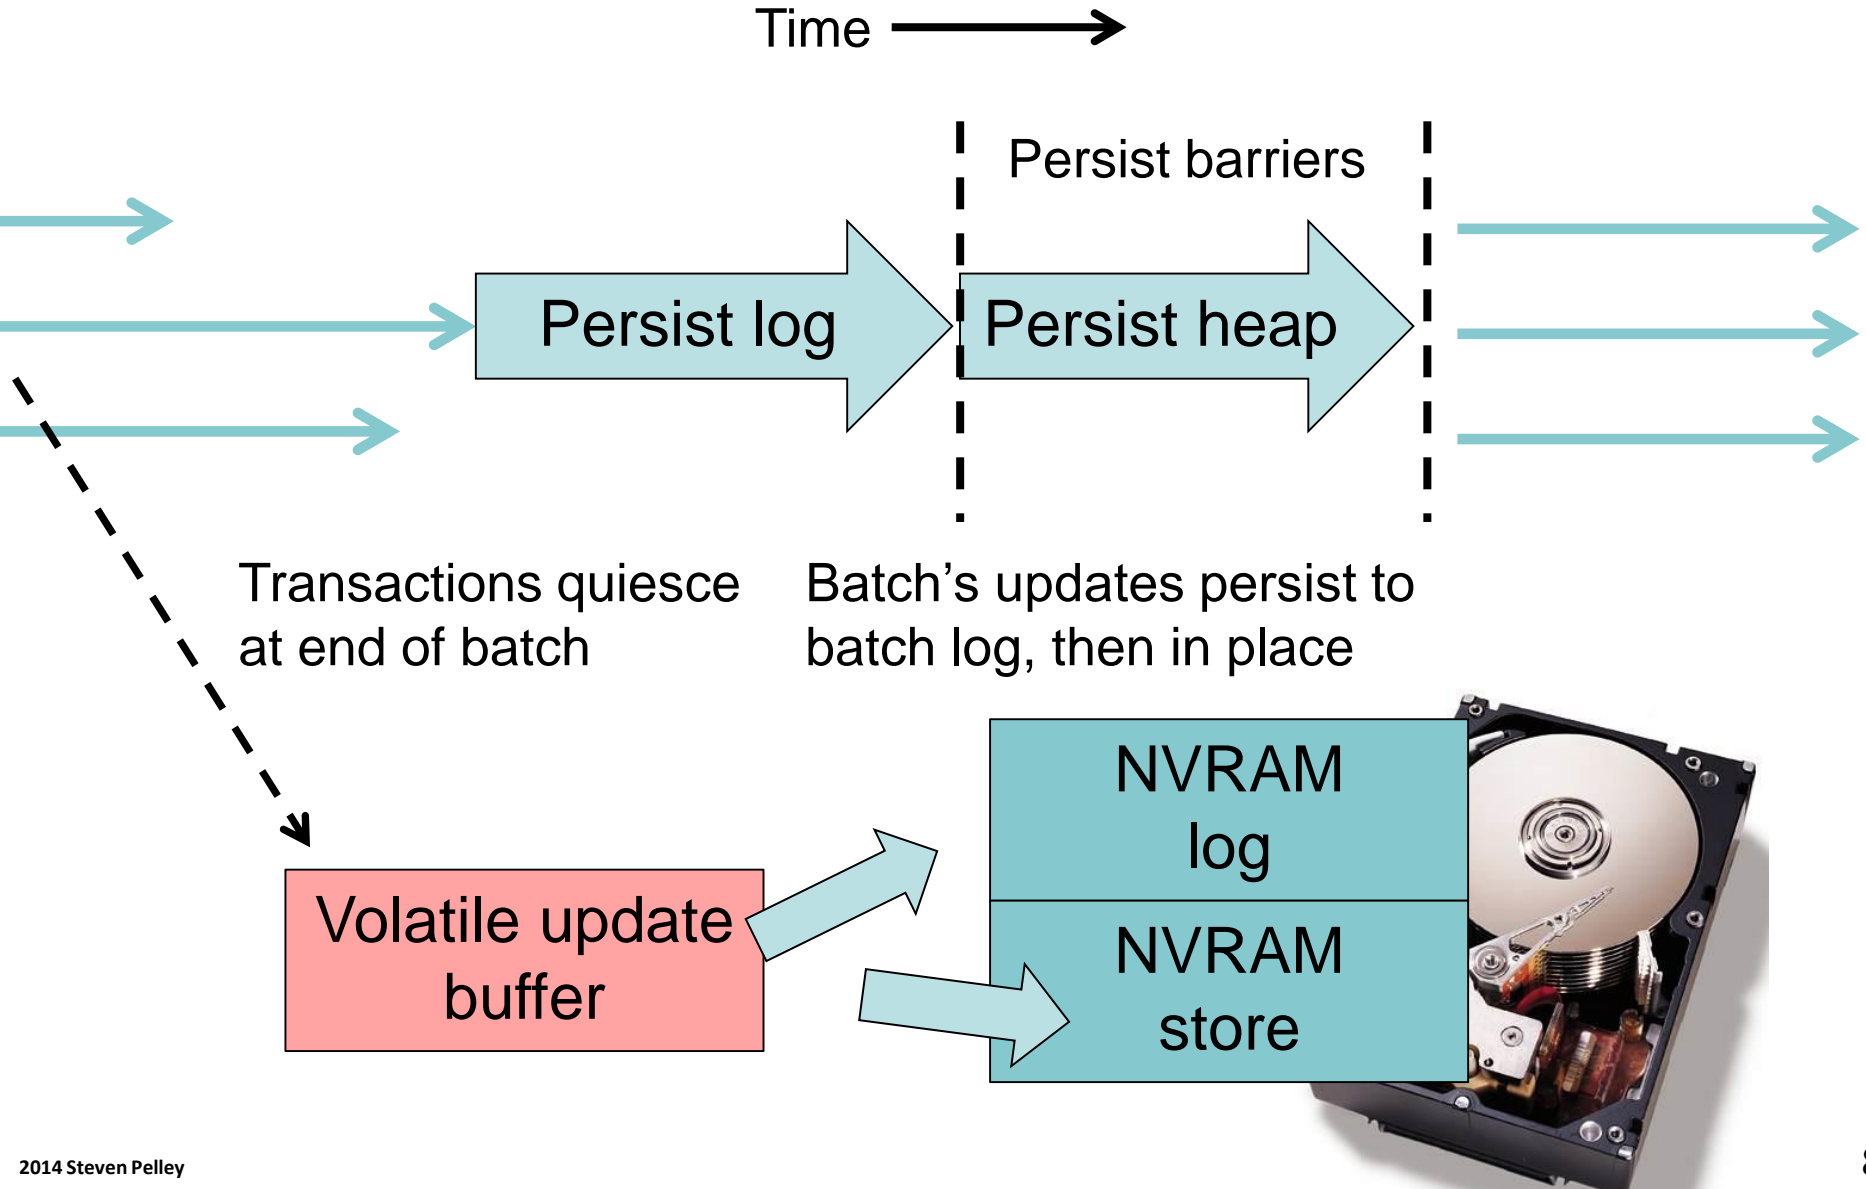
# NVRAM in-place updates

Persist immediately  
Remove centralized log  
Barriers necessary at each update



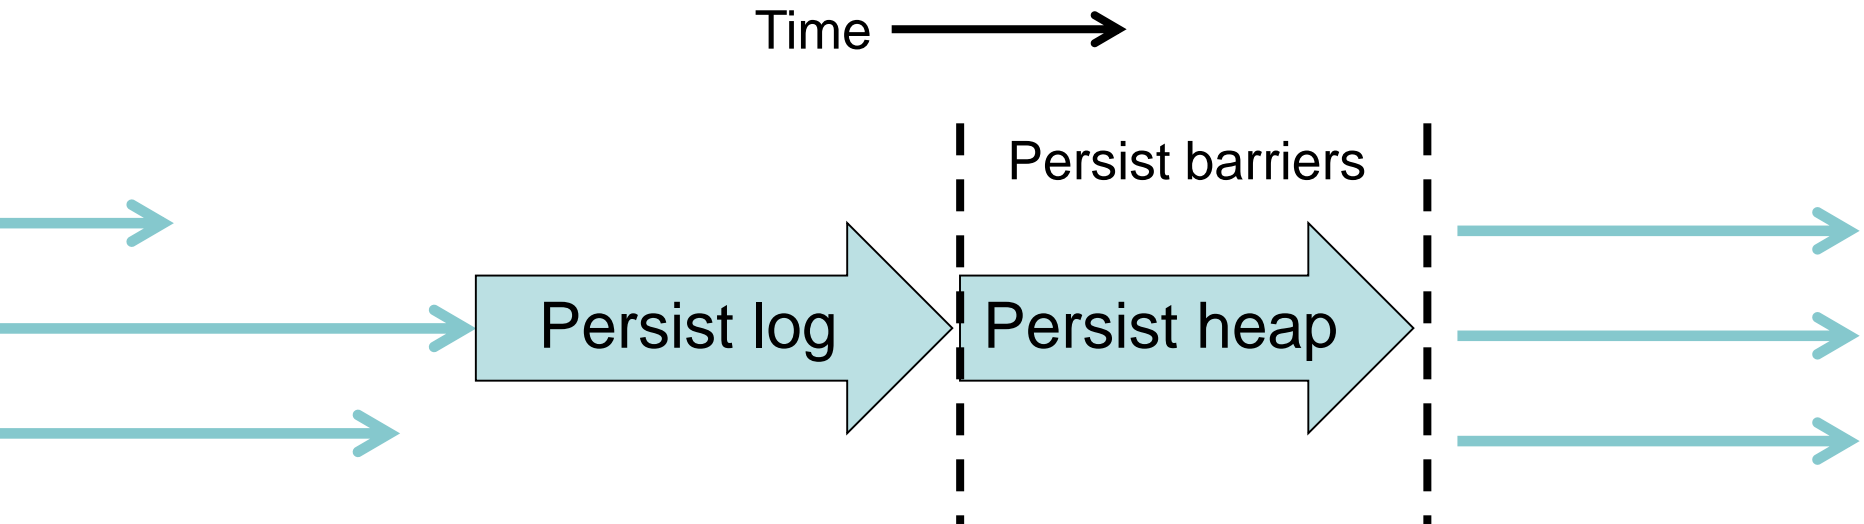
*Simple, removes software overhead, but frequent persist ordering*

# NVRAM Group Commit





# NVRAM Group Commit



*Constant number of persist barriers per batch*

*No centralized logging*


*Throughput high so long as batch quiesce/persist short*

*Several other tradeoffs (e.g., increased transaction latency)*


# Modeling unavailable devices

- Run database on *real hardware*
  - Log and db heap on RAMDisk (or just in DRAM)
  - Introduce precise delays (20ns precision using x86 RDTSCP) to model persist barrier latency
- Build recovery mechanisms in software
  - Shore-MT: research platform for high performance transaction processing
  - Rely on dirty bit fields to track buffer pool writes during transaction, page latch, or batch

# Recovery management performance




# Recovery management performance




*Disk replacement insensitive to latency, but slow*

# Recovery management performance



*Group commit recovers throughput when barrier latency exposed*

# Recovery management performance



*High performance barriers and memory system ideal*

## In the VLDB paper...

- Implementation and methodology details
- Persist bandwidth modeling
- Group commit trade-offs and transaction latency analysis
- Read-cache performance based on memory trace analysis

# Future work: NVRAM optimizations

- Asynchronous persists
  - Allow execution to proceed ahead of persists
  - Hides persist and persist barrier latency
  - *Persist ordering* critical path limits persist rate
- Persist coalescing
  - Omit persists if ordering constraints not violated
  - Acts as bandwidth filter (like write-back cache)
  - Reduces persist ordering critical path

*Need mechanisms to precisely label persist ordering constraints  
[upcoming ISCA: **Memory Persistency**]*



# Conclusion

- Disk-based software carries baggage
- Frequent persist synchronization also slow
- New software and memory systems improve performance and simplify software design
- References
  - Website: <http://web.eecs.umich.edu/~spelley>
  - Steven Pelley, Thomas F. Wenisch, Brian T. Gold, Bill Bridge: Storage Management in the NVRAM Era. PVLDB 7(2): 121-132 (2013)
  - [BPFS] Jeremy Condit, Edmund B. Nightingale, Christopher Frost, Engin Ipek, Benjamin Lee, Doug Burger, and Derrick Coetzee. 2009. Better I/O through byte-addressable, persistent memory. In *Proceedings of the ACM SIGOPS 22nd symposium on Operating systems principles (SOSP '09)*